



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

## High-Performance Computing Lab for CSE

2024

Student: Benedict Armstrong

Discussed with: Tristan Gabl

---

## Solution for Project 02

Due date: 25 March 2024, 23:59

---

**HPC Lab for CSE 2024 — Submission Instructions**  
(Please, notice that following instructions are mandatory:  
submissions that don't comply with, won't be considered)

- Assignments must be submitted to Moodle (i.e. in electronic format).
- Provide both executable package and sources (e.g. C/C++ files, Matlab). If you are using libraries, please add them in the file. Sources must be organized in directories called:  
*Project\_number\_lastname\_firstname*  
and the file must be called:  
*project\_number\_lastname\_firstname.zip*  
*project\_number\_lastname\_firstname.pdf*
- The TAs will grade your project by reviewing your project write-up, and looking at the implementation you attempted, and benchmarking your code's performance.
- You are allowed to discuss all questions with anyone you like; however: (i) your submission must list anyone you discussed problems with and (ii) you must write up your submission independently.

### 1. Computing $\pi$ with OpenMP [20 points]

### 2. The Mandelbrot set using OpenMP [20 points]

### 3. Bug hunt [10 points]

1. The first bug is a compile-time bug. The `#pragma` directive must be followed by a for loop. In this case we have a `tid = omp_get_thread_num()` statement immediately after the `#pragma`. This is not allowed. The fix is to move the `tid` assignment into the for loop.
2. There are a couple of errors in the code. The variable `tid` should be made explicitly private as every thread is writing to it. In the last for loop the total sum should be marked as a reduction variable (using the `reduction(+:total)` clause). Also by default the second loop will not spawn any new threads as the option `OMP_NESTED` is set to `FALSE` by default (see IBM OpenMP documentation).




Figure 1: Benchmark results for the  $\pi$  calculation using OpenMP.

#### 4. Parallel histogram calculation using OpenMP [15 points]

#### 5. Parallel loop dependencies with OpenMP [15 points]

To parallelize the loop with dependencies we split up the loop into  $N$  equal parts, where  $N$  is the number of threads. We then calculate the first element of each thread's partition

$$S_i = S_n * up^{i*chunk\_size}$$

where  $i$  is the thread number. Each thread then calculates the rest of the assigned elements.

#### 6. Quicksort using OpenMP tasks [20 points]

Quicksort can be easily parallelized using tasks. We create a task for each recursive call to the quicksort function. We then wait for all tasks to finish before returning. This is done by adding a `taskwait` directive after the recursive calls. The only complication is defining a minimum task size to prevent the creation of too many tasks. The code is shown in Listing 1.

After implementing the quicksort function using tasks I benchmarked the code using a strong scaling analysis. The results are shown in Figure 7.




Figure 2: Mandelbrot set

```
#pragma omp task shared(data) firstprivate(right) final(right < MIN_SIZE)  
quicksort(data, right);  
  
int t = length - left;  
#pragma omp task shared(data, left) firstprivate(t) final(t < MIN_SIZE)  
quicksort(&(data[left]), t);  
  
#pragma omp taskwait
```

Listing 1: Recursion of the quicksort function using tasks

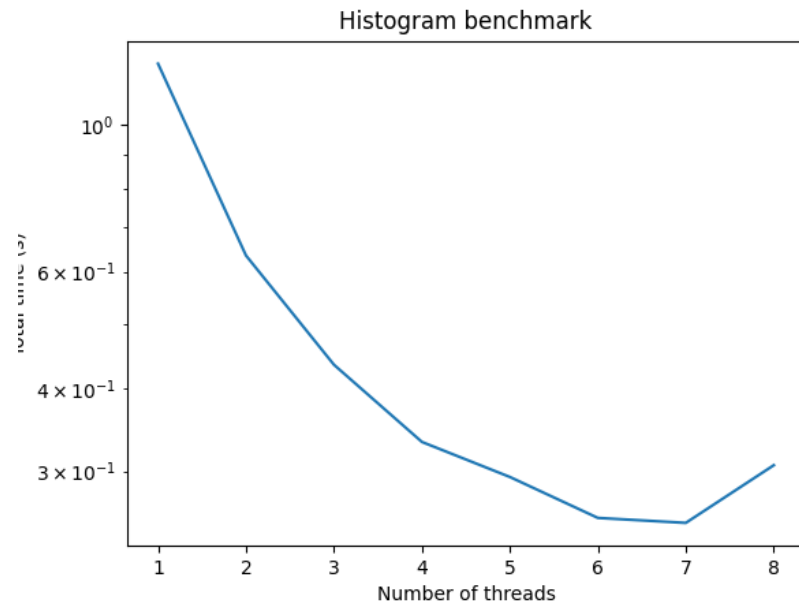


Figure 3: Benchmark results for the histogram calculation using OpenMP.




Figure 4: Speedup results for the histogram calculation using OpenMP.




Figure 5: Benchmark results for the loop dependencies calculation using `OpenMP`.

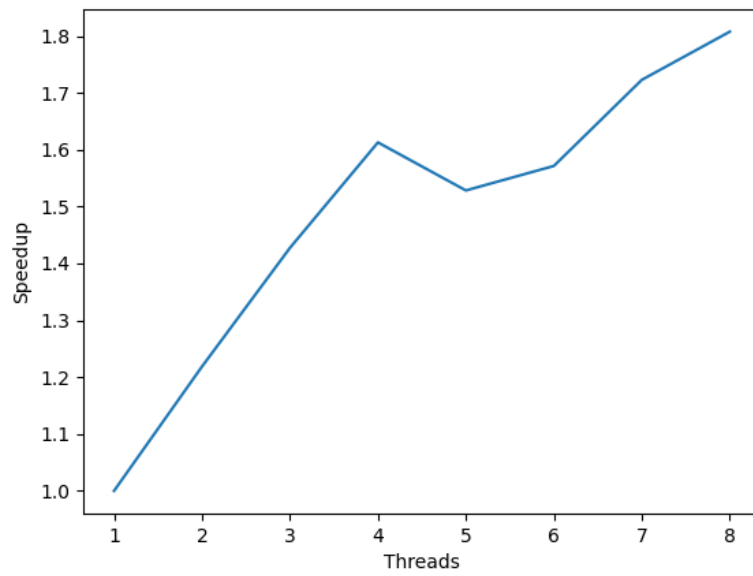


Figure 6: Speedup results for the loop dependencies calculation using `OpenMP`.

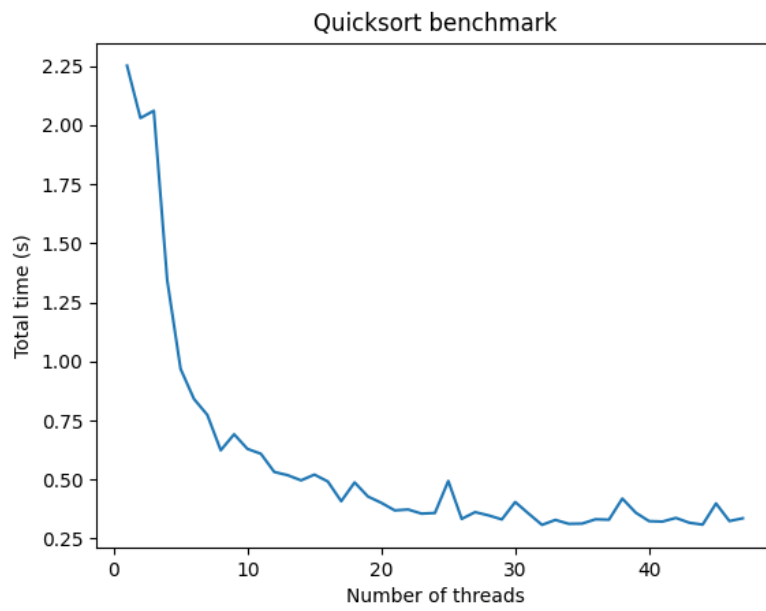


Figure 7: Benchmark results for the quicksort calculation using OpenMP tasks.